



---

# Distributed Hyperparameter Optimization and Model Search with Examples using SHADHO

---

A horizontal line composed of many small gold dots, spanning the width of the slide.

Jeffery Kinnison, Sadaf Ghaffari, Nathaniel Blanchard,  
Walter Scheirer



# Overview

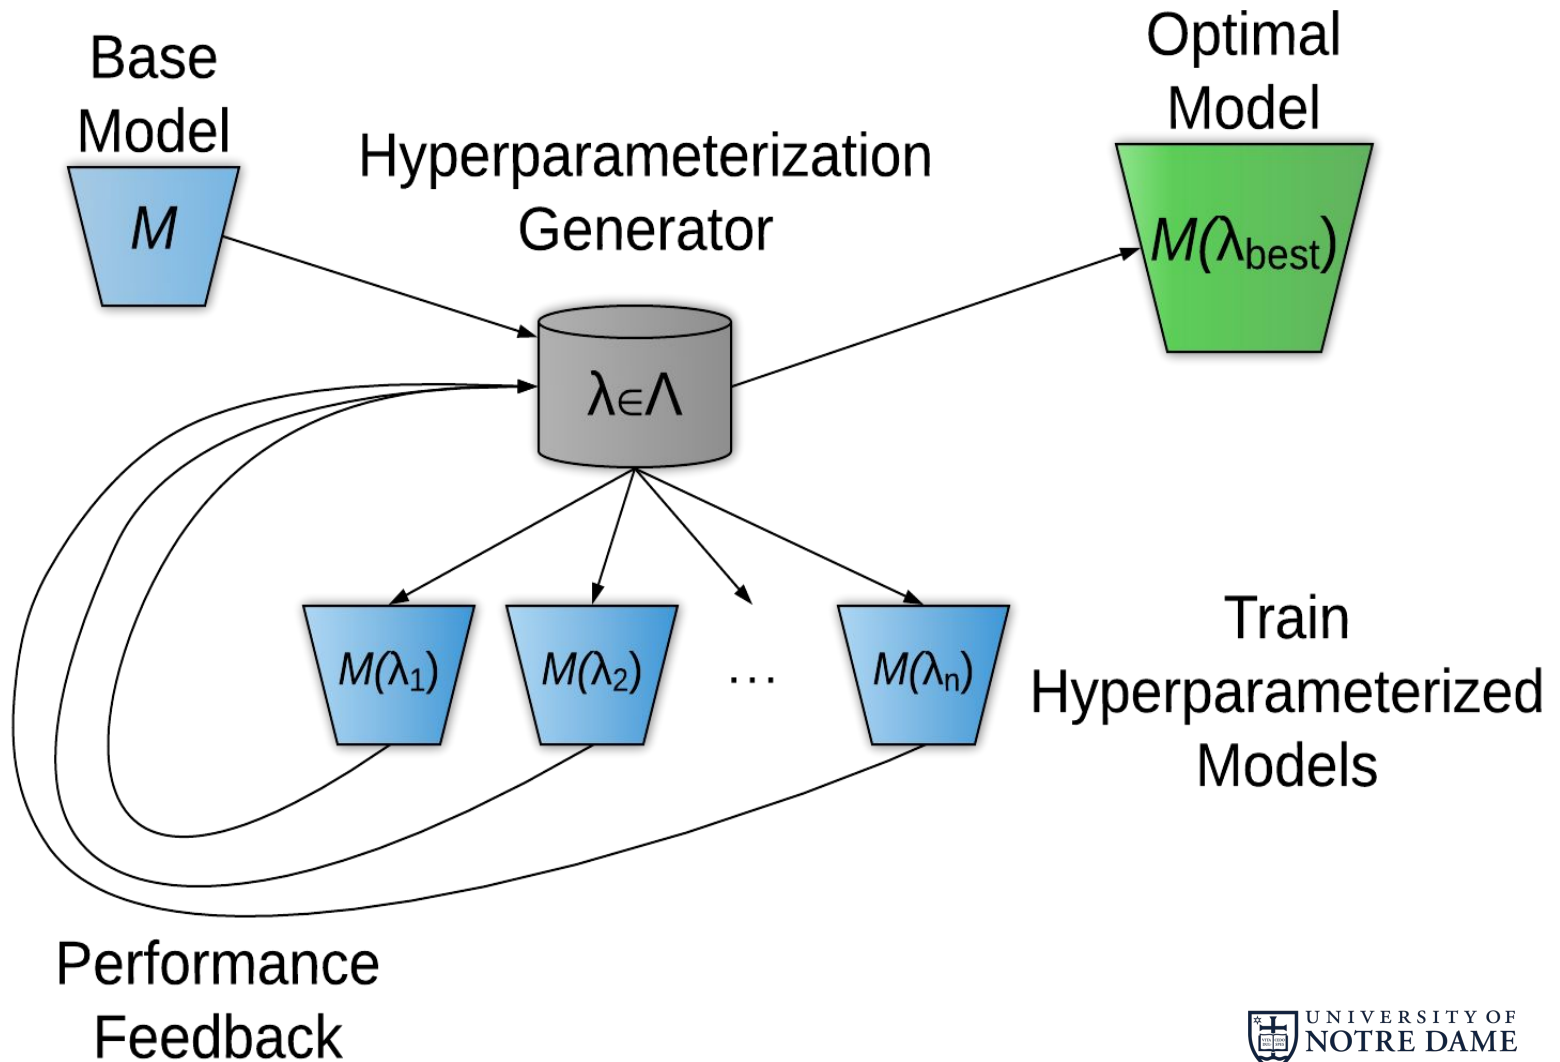
1. Scalable Hardware-Aware Distributed Hyperparameter Optimization
2. A First Example
3. Optimizing SVM
4. Optimizing Neural Network Structure
5. Going Forward



# 1. Scalable Hardware-Aware Distributed Hyperparameter Optimization

<https://jeffkinnison.github.io/shadho-tutorial/>


# Distributed Hyperparameter Optimization



# Structure of a Hyperparameter Search



CNN Architecture Search Domains




SVM Kernel Hyperparameter Domains

# Structure of a Hyperparameter Search

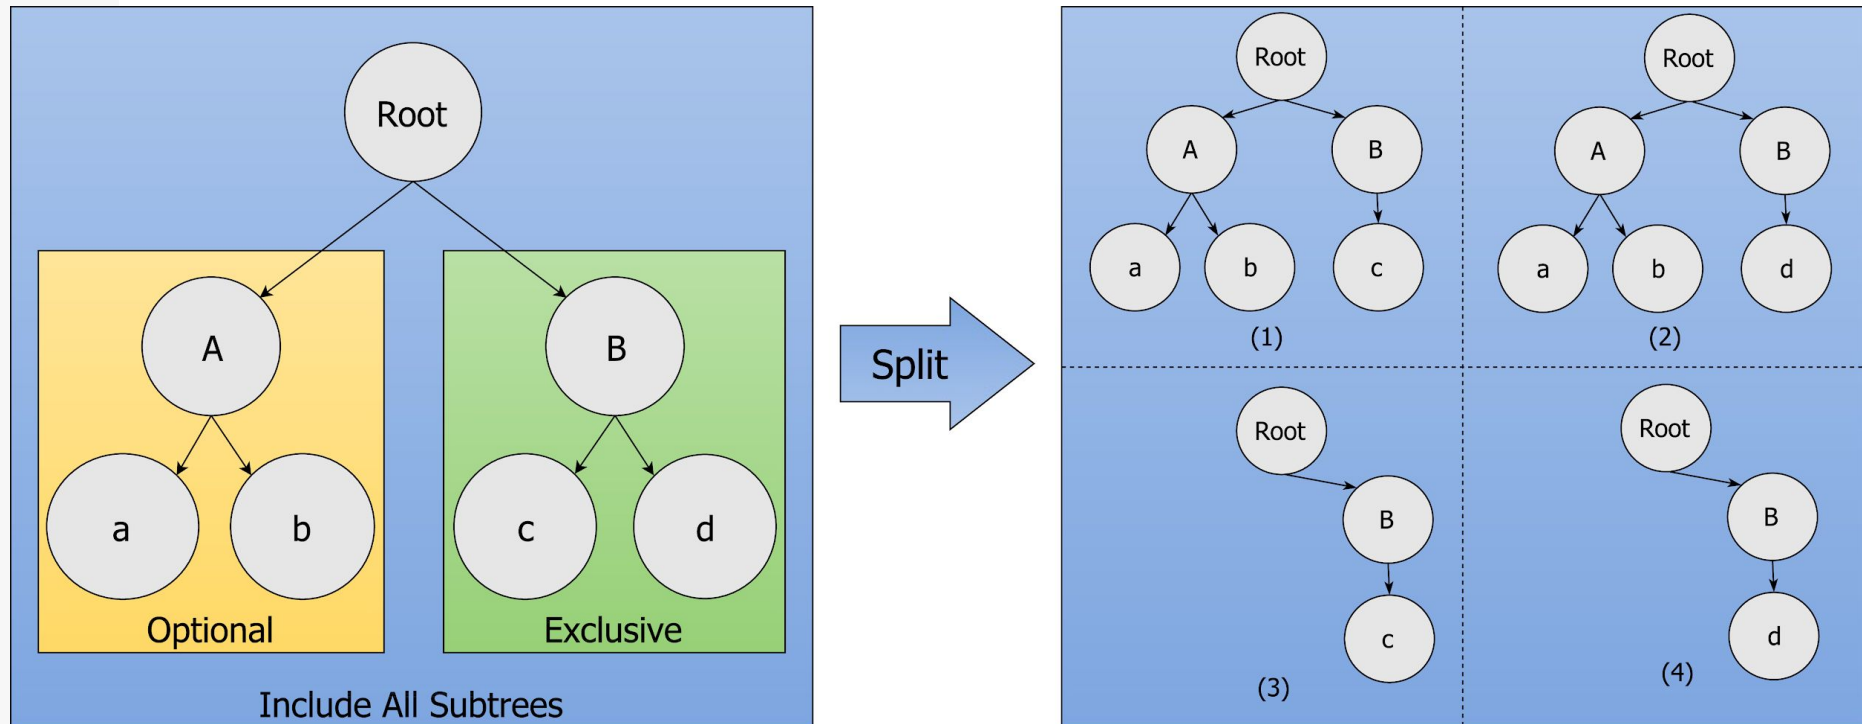


CNN Architecture Search Domains



SVM Kernel Hyperparameter Domains

# Properties of Search Spaces




# Heuristics for Scheduling: Complexity

Combinatorial size of hyperparameter search domains in a model.

Discrete Domain:

$$['a', 'b', 'c', 'd'] \Rightarrow 4$$

Continuous Domain:

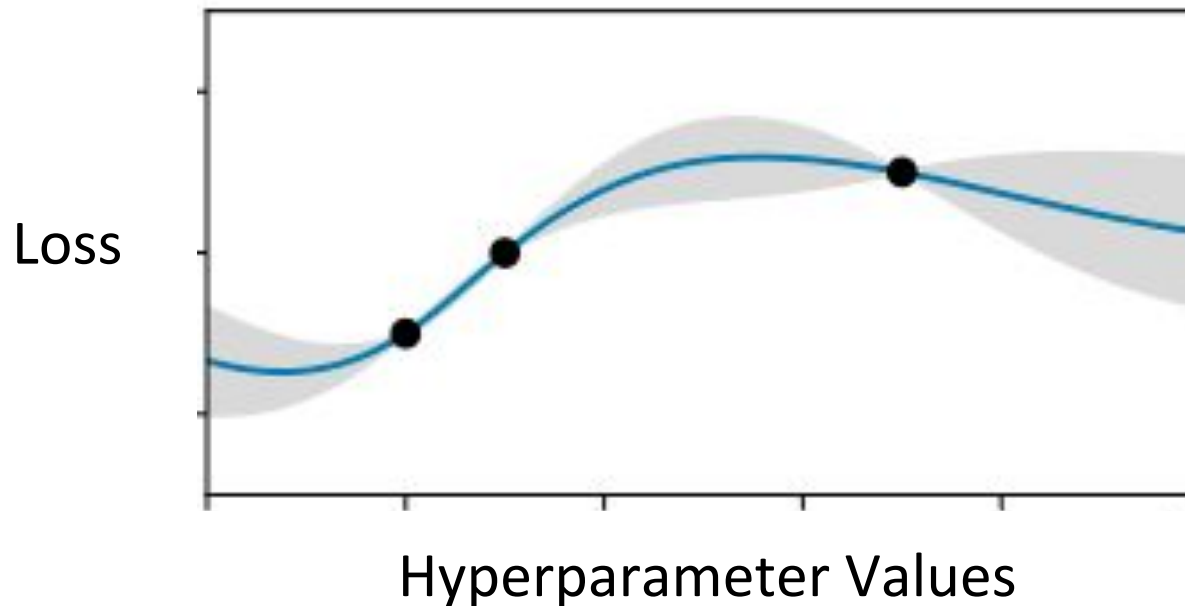


$$\Rightarrow 5.15$$



# Heuristics for Scheduling: Uncertainty

Approximation of confidence in model performance across different hyperparameterizations.



# Compute Classes

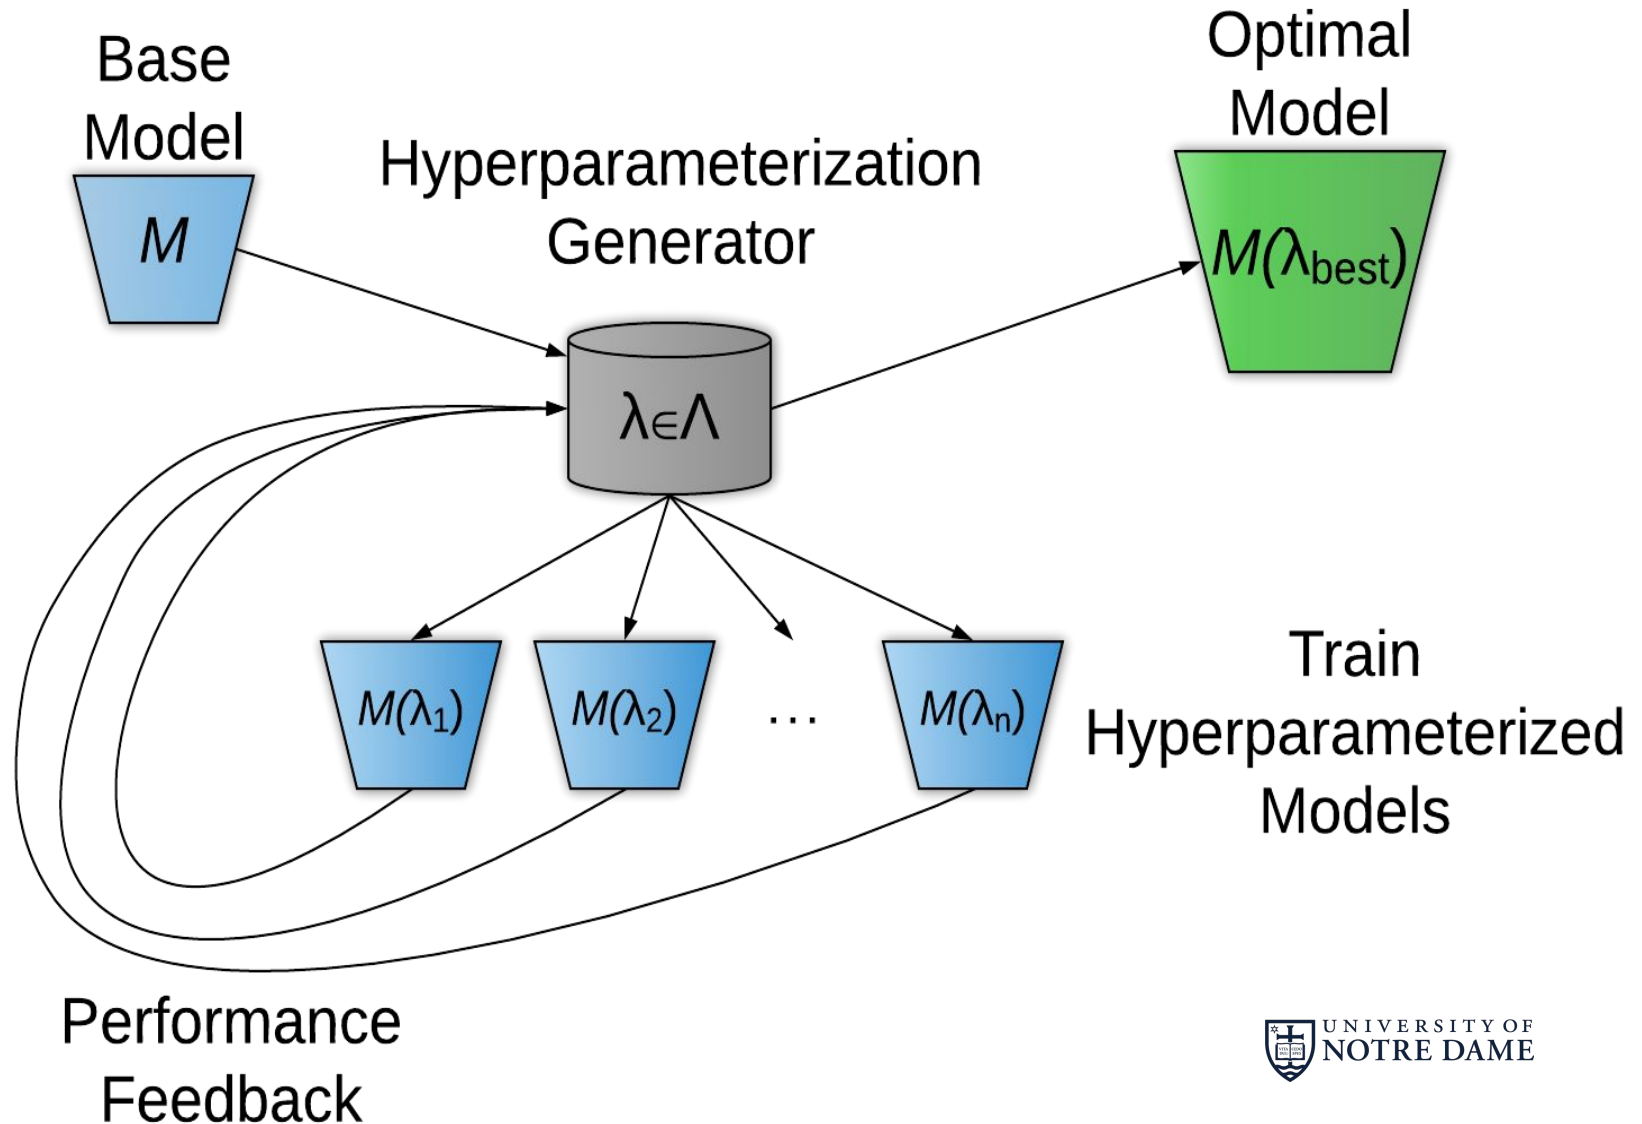
Observation: Different hardware offers different running time performance.

Strategy: Group compute nodes with common hardware together for predictable performance.

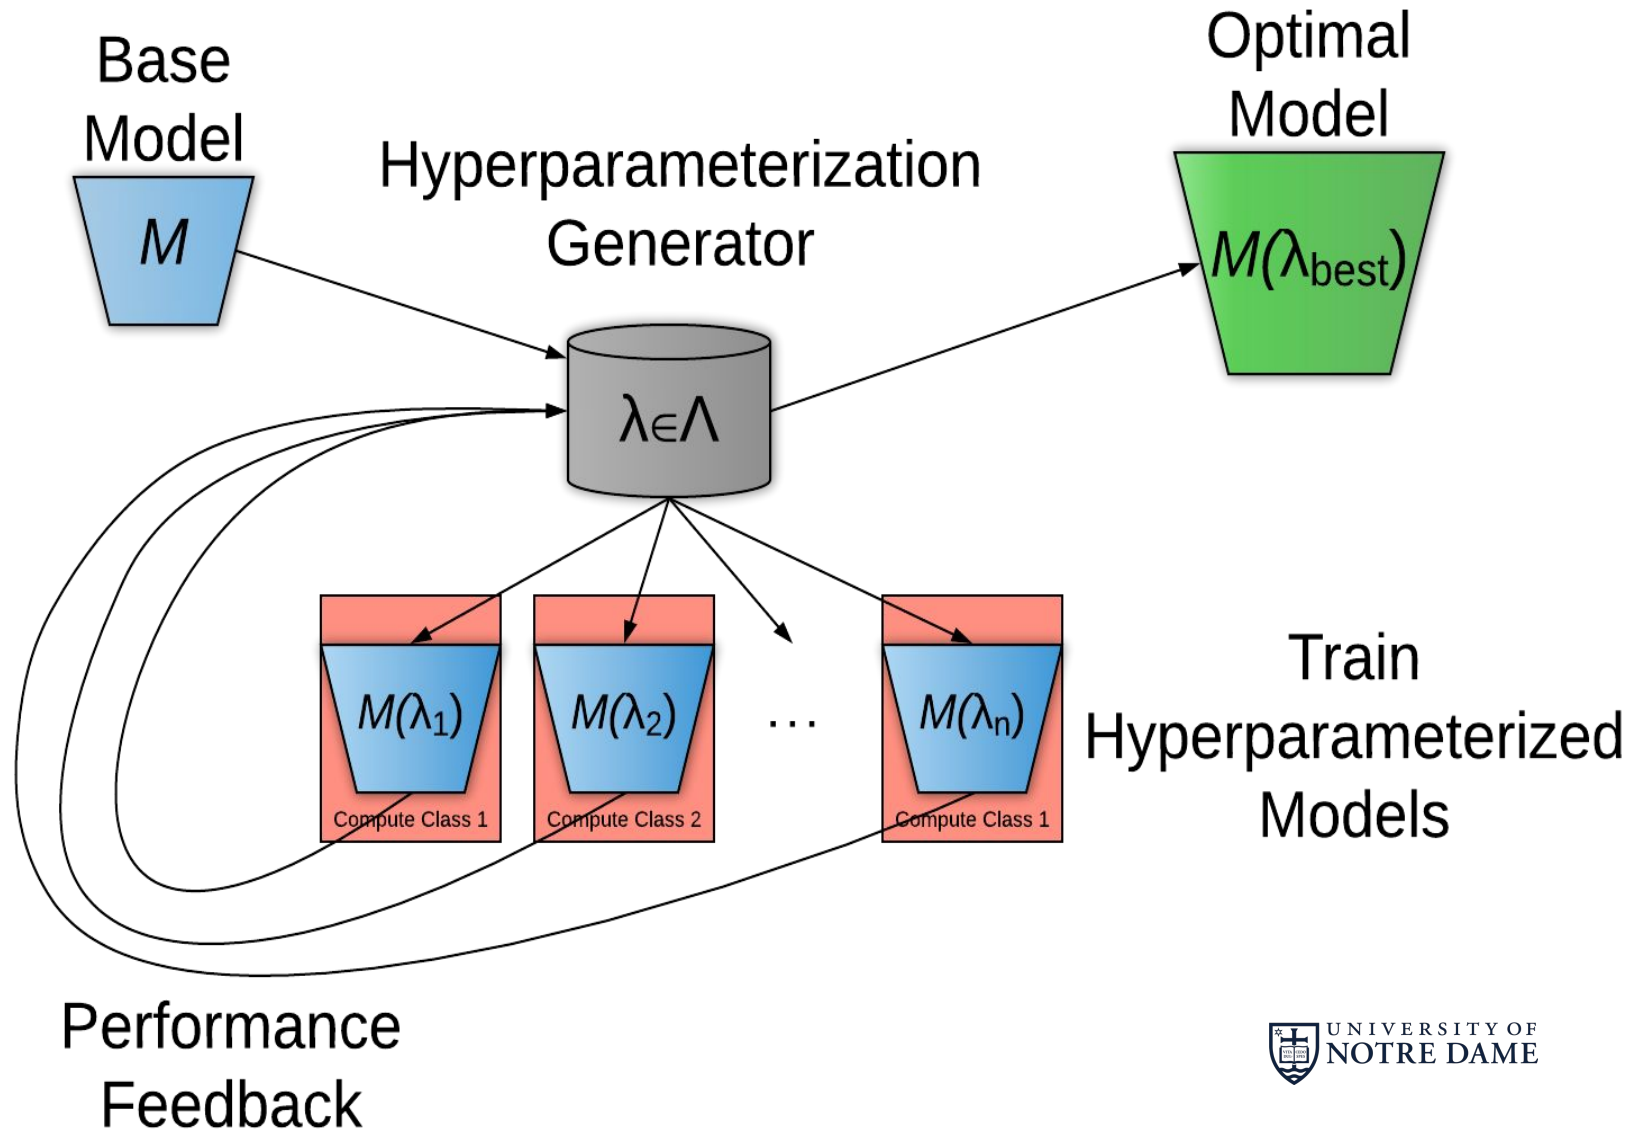
# SHADHO Optimzation

1. Organize hyperparameter searches into a forest of non-overlapping search spaces
2. Order trees in the forest by their structure (search **complexity**) and performance variability (search **uncertainty**)
3. Assign trees to **compute classes**, with high-complexity, high-priority trees assigned to higher-performing hardware
4. Evaluate parameterized models on assigned hardware
5. Update priority and repeat

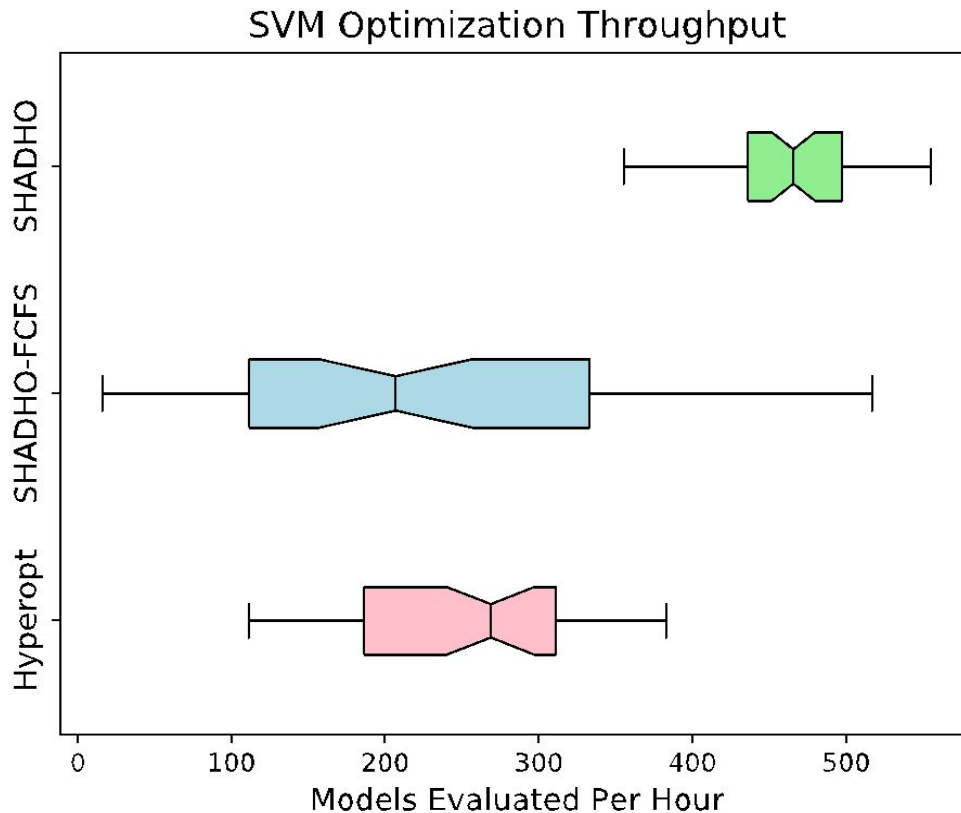
# Distributed Hyperparameter Optimization



# SHADHO Optimization



# SVM Kernel Optimization for MNIST



On average over 48 trials, SHADHO increased search throughput by 1.8x.

# Defining a Search Space

## Continuous Domains

Draw values from a continuous probability distribution.

## Discrete/Categorical Domains

Randomly sample from a set of values.

## Grid Domains

Exhaustively search over a set of values.

## Dependent Domains

Compute a value in response to another hyperparameter.

# Defining a Search Space

## Hierarchical Domains

Organize domains hierarchically into trees.

## Repeating Domains

Create a repeating structure and search over subsets.

## Sequential Domains

Set up ordered sequences of domains.



# Available Sampling Methods

## Grid Search

Exhaustively sample over a regular grid.

## Random Search

Randomly sample with no guidance.

## Bayesian Optimization

- Gaussian Processes
- Random Forests
- TPE

# Getting SHADHO

<https://github.com/jeffkinnison/shadho>

```
pip3 install shadho  
  
python3 -m shadho.install.workqueue
```



## 2. A First Example

<https://jeffkinnison.github.io/shadho-tutorial/>

# Optimizing a Convex Function

Example 1:  $\sin(x) * \cos(y)$

In this example, we will:

- Create continuous and categorical search domains
- Set up a SHADHO driver
- Run the search locally

# Optimizing a Convex Function

Example 1: Creating the search space

```
from shadho import Shadho, spaces

search_space = {
    'x': spaces.uniform(0, 2 * math.pi),
    'y': spaces.choice(list(np.linspace(0, 2 * math.pi, 1000)))
}
```

## Search Domains

(Continuous)  $x \sim U(0, 2\pi)$

(Categorical)  $y \sim$  grid of 1000 points in  $[0, 2\pi]$

# Optimizing a Convex Function

Example 1: Setting up the objective

```
def objective(params):  
    x = params['x']  
    y = params['y']  
  
    return np.sin(x) * np.cos(y)
```

Receives a dictionary of parameters

Returns a floating-point result

# Optimizing a Convex Function

## Example 1: Setting up the driver

```
opt = Shadho(  
    'convex-tutorial', # Name of this experiment  
    objective,         # The function to optimize  
    search_space,      # The search space to sample  
    method='random',   # The sampling method  
    timeout=30         # The time to run the search (s).  
)
```

Pass the objective function to the driver with the search space.

# Optimizing a Convex Function

Example 1: Setting up the driver

```
opt = Shadho(  
    'convex-tutorial', # Name of this experiment  
    objective,         # The function to optimize  
    search_space,      # The search space to sample  
    method='random',   # The sampling method  
    timeout=30         # The time to run the search (s).  
)  
  
opt.run()
```



# Optimizing a Convex Function

Example 1: Running the search

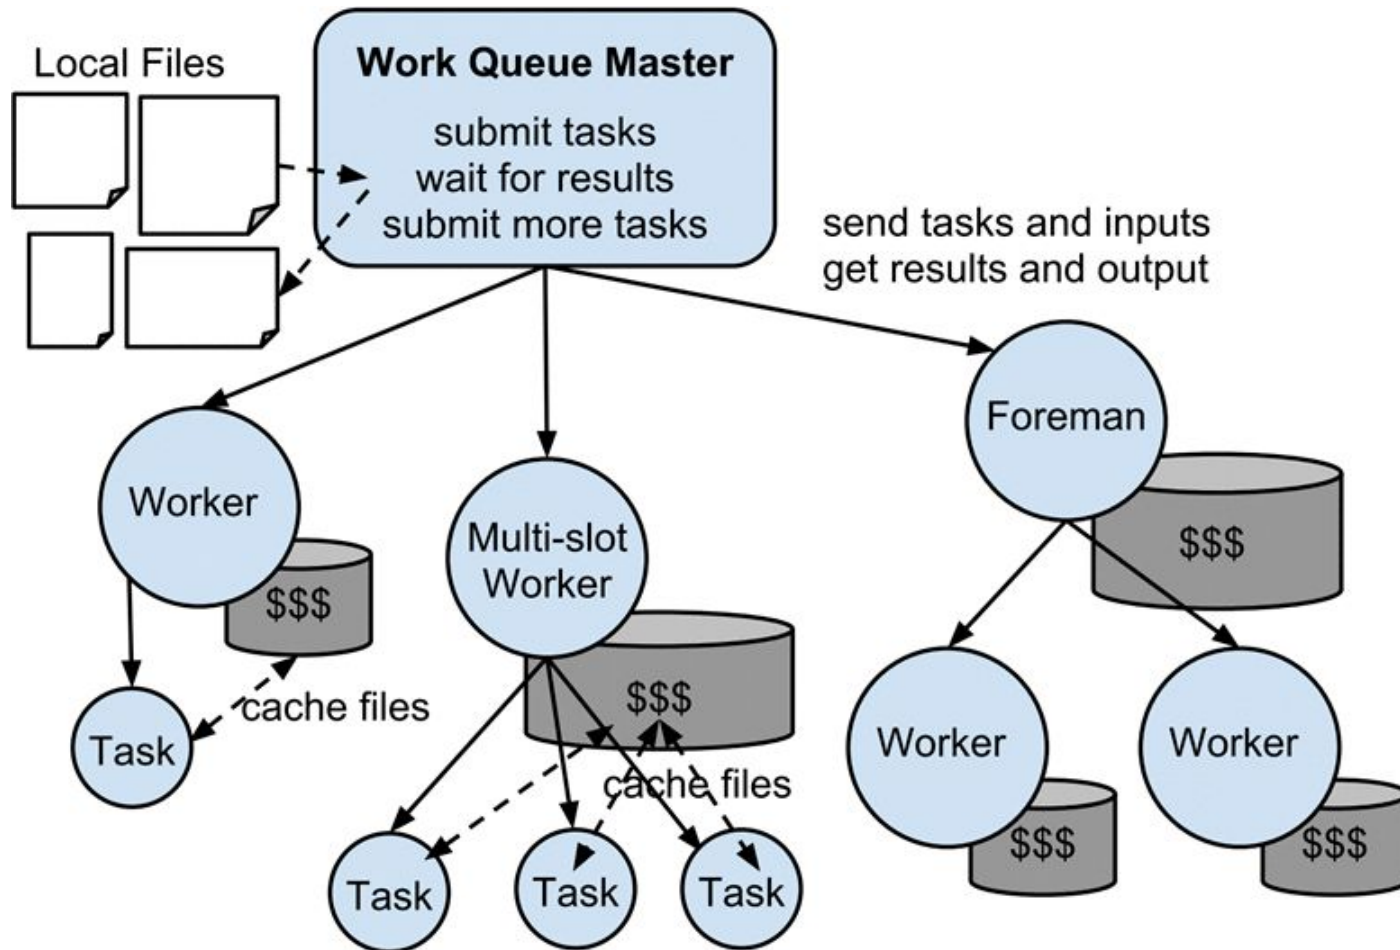
```
opt = Shadho(  
    'convex-tutorial', # Name of this experiment  
    objective,        # The function to optimize  
    search_space,      # The search space to sample  
    method='random',   # The sampling method  
    timeout=30         # The time to run the search (s).  
)  
  
opt.run()
```



# 3. Distributed Search with SHADHO

<https://jeffkinnison.github.io/shadho-tutorial/>

# Distributing Trials with Work Queue



# Distributing Trials with Work Queue

- Master/worker model -> workers ping master for tasks
- Master exists as a public-facing server
- Workers can connect from any domain
- Workers advertise hardware capabilities of host machine
- Workers operate in a sandbox, cache persistent files
- Communicates in files sent to worker/returned to master

# Distributing Trials

Example 2:  $\sin(x) * \cos(y)$

In this example, we will:

- Convert the previous example to distributed
- Demonstrate running a SHADHO worker

# Optimizing a Convex Function

Example 2: Creating the search space

```
from shadho import Shadho, spaces

search_space = {
    'x': spaces.uniform(0, 2 * math.pi),
    'y': spaces.choice(list(np.linspace(0, 2 * math.pi, 1000)))
}
```

## Search Domains

(Continuous)  $x \sim U(0, 2\pi)$

(Categorical)  $y \sim$  grid of 1000 points in  $[0, 2\pi]$

# Optimizing a Convex Function

## Example 2: Setting up the driver

```
opt = Shadho(  
    'distributed-tutorial', # Name of this experiment  
    'bash evaluate.sh',    # The function to optimize  
    search_space,          # The search space to sample  
    method='random',       # The sampling method  
    timeout=30             # The time to run the search (s)  
)
```

Pass a command to be run on each worker to the driver with the search space.

# Optimizing a Convex Function

Example 2: Setting up the objective (*objective.py*)

```
import math

def objective(params):
    x = params['x']
    y = params['y']
    return math.sin(x) * math.cos(y)

if __name__ == '__main__':
    import shadho_worker
    shadho_worker.run(objective)
```

The objective function is moved to its own file.

Wrap the objective with `shadho_worker`, which handles parsing hyperparameters.



# Optimizing a Convex Function

Example 2: Setting up the objective (*evaluate.sh*)

```
#!/usr/bin/env bash  
  
# Load any environment to run the objective.  
  
python3 objective.py
```

The worker will run this file. Use it like any other shell script, job file, etc.

# Optimizing a Convex Function

Example 2: Adding *evaluate.sh* and *objective.py*

```
opt.add_input_file('evaluate.sh')  
opt.add_input_file('objective.py')
```

Give SHADHO the path to the input files.

These files will be cached on every worker between trials.

# Optimizing a Convex Function

Example 2: Running the driver and worker

Driver - In one terminal:

```
python3 driver.py
```

Worker - In another terminal:

```
python3 -m shadho.workers.workqueue \  
-M distributed-tutorial # experiment key
```



## 6. Optimizing SVM

<https://jeffkinnison.github.io/shadho-tutorial/>

# Optimizing SVM Kernels

Example 2: *SVM for MNIST Classification*

In this example, we will:

- Create hierarchical search spaces
- Work with compute classes
- Return additional results

# Optimizing SVM Kernels

Example 2: Setting up the search space

```
C = spaces.uniform(-1000, 2000)
gamma = spaces.log10_uniform(-5, 8)
coef0 = spaces.uniform(-1000, 2000)
degree = [2, 3, 4, 5, 6, 7]
```

Search domains can be stored and used multiple times.

# Optimizing SVM Kernels

## Example 2: Setting up the search space

```
search_space = {  
    'linear': {  
        'C': C,  
    },  
    'rbf': {  
        'C': C,  
        'gamma': gamma,  
    },  
    'sigmoid': {  
        'C': C,  
        'gamma': gamma,  
        'coef0': coef0,  
    },  
    'poly': {  
        'C': C,  
        'gamma': gamma,  
        'coef0': coef0,  
        'degree': degree,  
    },  
    'exclusive': True  
}
```

Each kernel is stored independently.

The “exclusive” flag tells SHADHO to sample each independently.

# Optimizing SVM Kernels

## Example 2: Setting up the driver

```
opt = Shadho(  
    'svm-tutorial', # The experiment key  
    'bash evaluate.sh', # The command to run on the worker  
    search_space, # The search space  
    method='random', # The sampling method to use  
    timeout=120 # The amount of time to run (s)  
)
```



# Optimizing SVM Kernels

## Example 2: Setting up the objective

```
from sklearn.svm import svc

def main(params):
    ... kernel = list(params.keys())[0]
    ... kernel_params = params[kernel]
    ...
    ...
    ... svc = SVC(kernel=kernel, **kernel_params)
```

# Optimizing SVM Kernels

## Example 2: Setting up the objective

```
def main(params):  
    ...  
    ...  
    out = {  
        ... 'loss': loss,  
        ... 'accuracy': accuracy,  
        ... 'precision': precision,  
        ... 'recall': recall,  
        ... 'train_time': train_time,  
        ... 'test_time': test_time  
    }  
    ...  
    return out  
  
if __name__ == '__main__':  
    ... import shadho_worker  
    ... shadho_worker.run(main)
```

Return multiple values in a dictionary and SHADHO will record them all.

“loss” is considered the objective.

# Optimizing a Convex Function

Example 2: Adding *evaluate.sh* and *train\_svm.py*

```
opt.add_input_file('evaluate.sh')  
opt.add_input_file('train_svm.py')  
opt.add_input_file('mnist.npz')
```

In addition to code, we can send small datasets to be cached on workers.

# Optimizing SVM Kernels

Example 2: Setting up the compute classes

```
opt.add_compute_class('16-core', 'cores', 16, max_tasks=20)  
opt.add_compute_class('8-core', 'cores', 8, max_tasks=25)  
opt.add_compute_class('4-core', 'cores', 4, max_tasks=50)
```

Compute classes are given a name, a resource grouping them, the amount of that resource, and the maximum expected number of workers.

# Optimizing a SVM Kernels

Example 2: Running the driver and worker

Driver - In one terminal:

```
python3 driver.py
```

Worker - In another terminal:

```
python3 -m shadho.workers.workqueue \  
.....-M svm-tutorial \  
.....--cores 4
```



# 5. Optimizing Neural Network Structure

<https://jeffkinnison.github.io/shadho-tutorial/>

# Optimizing Neural Network Structure

Example 2: *CNN for CIFAR-10 Classification*

In this example, we will:

- Demonstrate the object-oriented spaces API
- Use repeating and conditional domains
- Set up compute classes for GPU models

# Optimizing Neural Network Structure

Example 4: *CNN for CIFAR-10 Classification*

## Libraries Used:

- shadho
- pytorch
- torchvision (data & augmentation)
- ignite (training loop)



# Optimizing Neural Network Structure

## Example 2: Setting up the search space

```
activations = ['glu', 'leaky_relu', 'prelu', 'relu', 'selu', 'sigmoid', 'tanh']
batch_norm = spaces.log10_uniform(-4, 4)

conv_layer = spaces.scope(
    ··· out_filters=spaces.log2_randint(4, 10),
    ··· kernel_shape=spaces.randint(1, 10, step=2),
    ··· activation=activations,
    ··· batch_norm=batch_norm
)
```

# Optimizing Neural Network Structure

## Example 2: Setting up the search space

```
conv_layer.padding = spaces.dependent(  
    ... conv_layer.kernel_shape,  
    ... callback=lambda x: int(x // 2))
```

Dependent domains can be set up to react to other hyperparameter values.

This space implements “same” padding in response to generated kernel shapes

# Optimizing Neural Network Structure

Example 2: Setting up the search space

```
conv_layers = spaces.repeat(conv_layer, 6)
```

```
dense_layer = spaces.scope(  
    ... out_features=spaces.log2_randint(7, 13),  
    ... activation=activations,  
    ... batch_norm=batch_norm,  
)  
  
dense_layers = spaces.repeat(dense_layer, 3)
```

Domains may also be repeated  $n$  times to easily set up repeating structures.

# Optimizing Neural Network Structure

## Example 2: Building the network

```
layer_input = in_channels
for i, layer in enumerate(hyperparameters['conv_layers']):
    layers.append(nn.Conv2d(
        layer_input,
        layer['out_features'],
        layer['kernel_shape'],
        padding=layer['padding']))
    layers.append(activations[layer['activation']]())
    layers.append(nn.BatchNorm2d(layer['out_features']))
    layer_input = layer['out_features']

    if i % 2 == 1:
        layers.append(nn.MaxPool2d(2))
        data_shape = [int(d // 2) for d in data_shape]
```

# Optimizing Neural Network Structure

## Example 4: Building the network

```
layer_input = layer_input * np.product(data_shape)
layers.append(nn.Flatten())
for layer in hyperparameters['dense_layers']:
    layers.append(nn.Linear(layer_input, layer['out_features']))
    layers.append(activations[layer['activation']]())
    layer_input = layer['out_features']
```



## 6. Going Forward

<https://jeffkinnison.github.io/shadho-tutorial/>

# Adding to the Software

- More sampling methods (e.g., population-based)
- Additional hyperparameter domains
- Automatic compute class detection

# Updating the Process

- Explore uncertainty quantification and hyperparameter importance methods as heuristics
- Develop new search strategies
- Extending beyond hyperparameters





# THANKS!

---

**Any questions?**